

Markov Logic Networks

Andrea Passerini
passerini@disi.unitn.it

Statistical relational learning

Combining logic with probability

Motivation

- First-order logic is a powerful language to represent complex relational information
- Probability is the standard way to represent uncertainty in knowledge
- Combining the two would allow to model complex probabilistic relationships in the domain of interest

logic graphical models

- Graphical models are a mean to represent joint probabilities highlighting the relational structure among variables
- A compressed representation of such models can be obtained using templates, cliques in the graphs sharing common parameters (e.g. as in HMM for BN or CRF for MN)
- Logic can be seen as a language to build templates for graphical models
- Logic based versions of HMM, BN and MN have been defined

First-order logic (in a nutshell)

Symbols

- *Constant* symbols representing objects in the domain (e.g. `Nick`, `Polly`)
- *Variable* symbols which take objects in the domain as values (e.g. `x`, `y`)
- *Function* symbols which mapping tuples of objects to objects (e.g. `BandOf`). Each function symbol has an arity (i.e. number of arguments)
- *Predicate* symbols representing relations among objects or object attributes (e.g. `Singer`, `SangTogether`). Each predicate symbol has an arity.

Terms

- A *term* is an expression representing an object in the domain. It can be:
 - A constant (e.g. `Niel`)
 - A variable (e.g. `x`)
 - A function applied to a tuple of objects. E.g.:
`BandOf (Niel), SonOf (f, m), Age (MotherOf (John))`

Formulas

- A (well formed) *atomic formula* (or *atom*) is a predicate applied to a tuple of objects. E.g.:
`Singer(Nick), SangTogether(Nick, Polly)`
`Friends(X, BrotherOf(Emy))`
- Composite formulas are constructed from atomic formulas using logical connectives and quantifiers

First-order logic

Connectives

negation $\neg F$: true iff formula F is false

conjunction $F_1 \wedge F_2$: true iff both formulas F_1, F_2 are true

disjunction $F_1 \vee F_2$: true iff at least one of the two formulas F_1, F_2 is true

implication $F_1 \Rightarrow F_2$ true iff F_1 is false or F_2 is true (same as $F_2 \vee \neg F_1$)

equivalence $F_1 \Leftrightarrow F_2$ true iff F_1 and F_2 are both true or both false (same as $(F_1 \Rightarrow F_2) \wedge (F_2 \Rightarrow F_1)$)

Literals

- A *positive literal* is an atomic formula
- A *negative literal* is a negated atomic formula

Quantifiers

existential quantifier $\exists x F_1$: true iff F_1 is true for at least one object x in the domain. E.g.:

$\exists x \text{ Friends}(x, \text{BrotherOf}(\text{Emy}))$

universal quantifier $\forall x F_1$: true iff F_1 is true for all objects x in the domain. E.g.:

$\forall x \text{ Friends}(x, \text{BrotherOf}(\text{Emy}))$

Scope

- The *scope* of a quantifier in a certain formula is the (sub)formula to which the quantifiers applies

First-order logic

Precedence

- Quantifiers have the highest precedence
- Negation has higher precedence than other connectives
- Conjunction has higher precedence than disjunction
- Disjunction have higher precedence than implication and equivalence
- Precedence rules can as usual be overruled using parentheses

Examples

- *Emy and her brother have no common friends:*

$\neg \exists x (\text{Friends}(x, \text{Emy}) \wedge \text{Friends}(x, \text{BrotherOf}(\text{Emy})))$

- *All birds fly:*

$\forall x (\text{Bird}(x) \Rightarrow \text{Flies}(x))$

Closed formulas

- A variable-occurrence within the scope of a quantifier is called *bound*. E.g. x in:

$$\forall x (\text{Bird}(x) \Rightarrow \text{Flies}(x))$$

- A variable-occurrence outside the scope of any quantifier is called *free*. e.g. y in:

$$\neg \exists x (\text{Friends}(x, \text{Emy}) \wedge \text{Friends}(x, y))$$

- A *closed* formula is a formula which contains no free occurrence of variables

Note

- We will be interested in closed formulas only

Ground terms and formulas

- A *ground* term is a term containing no variables
- A *ground* formula is a formula made of only ground terms

First-order logic

First order language

- The set of symbols (constants, variables, functions, predicates, connectives, quantifiers) constitute a first-order *alphabet*
- A first order *language* given by the alphabet is the set of formulas which can be constructed from symbols in the alphabet


Knowledge base (KB)

- A first-order knowledge base is a set of formulas
- Formulas in the KB are implicitly conjoined
- A KB can thus be seen as a single large formula

Interpretation

- An *interpretation* provides *semantics* to a first order language by:
 - 1 defining a domain containing all possible objects
 - 2 mapping each ground term to an object in the domain
 - 3 assigning a truth value to each ground atomic formula (a *possible world*)
- The truth value of complex formulas can be obtained combining interpretation assignments with connective and quantifier rules

First-order logic: example




First-order logic: example

$\neg \exists x (\text{Friends}(x, \text{Emy}) \wedge \neg \text{Friends}(x, \text{BrotherOf}(\text{Emy})))$

- The formula is true under the interpretation as the following atomic formulas are true:

`Friends(Ann, Emy)`
`Friends(Ann, BrotherOf(Emy))`
`Friends(Matt, Emy)`
`Friends(Matt, BrotherOf(Emy))`
`Friends(John, Emy)`
`Friends(John, BrotherOf(Emy))`



Types

- Objects can be *typed* (e.g. people, cities, animals)
- A typed variable can only range over objects of the corresponding type
- a typed term can only take arguments from the corresponding type. E.g.

`MotherOf(John)`, `MotherOf(Amy)`

First-order logic

Inference in first-order logic

- A formula F is *satisfiable* iff there exists an interpretation under which the formula is true
- A formula F is *entailed* by a KB iff it is true for all interpretations for which the KB is true. We write it:

$$KB \models F$$

the formula is a logical consequence of KB, not depending on the particular interpretation

- Logical entailment is usually done by *refutation*: proving that $KB \wedge \neg F$ is unsatisfiable

Note

- Logical entailment allows to *extend* a KB inferring new formulas which are true for the same interpretations for which the KB is true

Clausal form

- The *clausal* form or *conjunctive normal form* (CNF) is a regular form to represent formulas which is convenient for automated inference:
 - A *clause* is a disjunction of literals.
 - A KB in CNF is a conjunction of clauses.
- Variables in KB in CNF are always implicitly assumed to be universally quantified.
- Any KB can be converted in CNF by a mechanical sequence of steps
- Existential quantifiers are replaced by Skolem constants or functions

Conversion to clausal form: example

First Order Logic	Clausal Form
<p>“Every bird flies”</p> $\forall x (\text{Bird}(x) \Rightarrow \text{Flies}(x))$	$\text{Flies}(x) \vee \neg \text{Bird}(x)$
<p>“Every predator of a bird is a bird”</p> $\forall x, y (\text{Predates}(x, y) \wedge \text{Bird}(y) \Rightarrow \text{Bird}(x))$	$\text{Bird}(x) \vee \neg \text{Bird}(y) \vee \neg \text{Predates}(x, y)$
<p>“Every prey has a predator”</p> $\forall y (\text{Prey}(y) \Rightarrow \exists x \text{Predates}(x, y))$	$\text{Predates}(\text{PredatorOf}(y), y) \vee \neg \text{Prey}(y)$

Problem of uncertainty

- In most real world scenarios, logic formulas are *typically* but not *always* true
- For instance:
 - “Every bird flies” : what about an ostrich (or Charlie Parker) ?
 - “Every predator of a bird is a bird”: what about lions with ostriches (or heroin with Parker) ?
 - “Every prey has a predator”: predators can be extinct
- A world failing to satisfy even a single formula would not be possible
- there could be no possible world satisfying all formulas

Handling uncertainty

- We can relax the hard constraint assumption on satisfying all formulas
- A possible world not satisfying a certain formula will simply be less likely
- The more formula a possible world satisfies, the more likely it is
- Each formula can have a weight indicating how strong a constraint it should be for possible worlds
- Higher weight indicates higher probability of a world satisfying the formula wrt one not satisfying it


Definition

- A Markov Logic Network (MLN) L is a set of pairs (F_i, w_i) where:
 - F_i is a formula in first-order logic
 - w_i is a real number (the weight of the formula)
- Applied to a finite set of constants $C = \{c_1, \dots, c_{|C|}\}$ it defines a Markov network $M_{L,C}$:
 - $M_{L,C}$ has one binary node for each possible grounding of each atom in L . The value of the node is 1 if the ground atom is true, 0 otherwise.
 - $M_{L,C}$ has one feature for each possible grounding of each formula F_i in L . The value of the feature is 1 if the ground formula is true, 0 otherwise. The weight of the feature is the weight w_i of the corresponding formula

Intuition

- A MLN is a *template* for Markov Networks, based on logical descriptions
- Single atoms in the template will generate nodes in the network
- Formulas in the template will generate cliques in the network
- There is an edge between two nodes iff the corresponding ground atoms appear together in at least one grounding of a formula in L

Markov Logic networks: example



Ground network

- A MLN with two (weighted) formulas:

$$w_1 \quad \forall x \text{ (Bird}(x) \Rightarrow \text{Flies}(x) \text{)}$$

$$w_2 \quad \forall x, y \text{ (Predates}(x, y) \wedge \text{Bird}(y) \Rightarrow \text{Bird}(x) \text{)}$$

- applied to a set of two constants `{Sparrow, Eagle}`
- generates the Markov Network shown in figure

Joint probability

- A ground MLN specifies a joint probability distribution over possible worlds (i.e. truth value assignments to all ground atoms)
- The probability of a possible world x is:

$$p(x) = \frac{1}{Z} \exp \left(\sum_{i=1}^F w_i n_i(x) \right)$$

where:

- the sum ranges over formulas in the MLN (i.e. clique templates in the Markov Network)
- $n_i(x)$ is the number of true groundings of formula F_i in x
- The partition function Z sums over all possible worlds (i.e. all possible combination of truth assignments to ground atoms)


Adding evidence

- Evidence is usually available for a subset of the ground atoms (as their truth value assignment)
- The MLN can be used to compute the conditional probability of a possible world x (consistent with the evidence) given evidence e :

$$p(x|e) = \frac{1}{Z(e)} \exp \left(\sum_{i=1}^F w_i n_i(x) \right)$$

- where the partition function $Z(e)$ sums over all possible worlds consistent with the evidence.

Example: evidence




Including evidence

- Suppose that we have (true) evidence e given by these two facts:

`Bird(Sparrow)`

`predates(Eagle, Sparrow)`

Example: assignment 1



Computing probability


$$p(x) = \frac{1}{Z} \exp(w_1 + 3w_2)$$

- The probability of a world with only evidence atoms set as true violates two ground formulas:

$\text{Bird}(\text{Sparrow}) \Rightarrow \text{Flies}(\text{Sparrow})$

$\text{Predates}(\text{Eagle}, \text{Sparrow}) \wedge \text{Bird}(\text{Sparrow}) \Rightarrow \text{Bird}(\text{Eagle})$

Example: assignment 2




Computing probability

$$p(x) = \frac{1}{Z} \exp(2w_1 + 4w_2)$$

- This possible world is the most likely among all possible worlds as it satisfies all constraints.

Example: assignment 3



Computing probability

$$p(x) = \frac{1}{Z} \exp(2w_1 + 4w_2)$$

- This possible world has also highest probability.
- The problem is that we did not encode constraints saying that:
 - A bird is not likely to be predator of itself
 - A prey is not likely to be predator of its predator

Hard constraints

Impossible worlds

- It is always possible to make certain worlds impossible by adding constraints with infinite weight
- Infinite weight constraints behave like pure logic formulas: any possible world has to satisfy them, otherwise it receives zero probability

Example


- Let's add the infinite weight constraint:

“Nobody can be a self-predator”

$w_3 \quad \forall X \neg \text{Predates}(X, X)$

to the previous example

Hard constraint: assignment 3




Computing joint probability

$$p(x) = \frac{1}{Z} \exp(2w_1 + 4w_2) = 0$$

- The numerator does not contain w_3 , as the no-self-predator constraint is never satisfied
- However the partition function Z sums over all possible worlds, including those in which the constraint is satisfied.
- As $w_3 = \infty$, the partition function takes infinite value and the possible worlds gets zero probability.

Hard constraint: assignment 2



Computing joint probability

$$p(x) = \frac{1}{Z} \exp(2w_1 + 4w_2 + 2w_3) \neq 0$$

- The only non-zero probability possible worlds are those always satisfying hard constraints
- Infinite weight features cancel out between numerator and possible worlds at denominator which also satisfy the constraints, while those which do not become zero

Assumptions

- For simplicity of presentation, we will consider MLN in form:
 - function-free (only predicates)
 - clausal
- However the methods can be applied to other forms as well
- We will use general first-order logic form when describing applications

MPE inference

- One of the basic tasks consists of predicting the most probable state of the world given some evidence (the *most probable explanation*)
- The problem is a special case of MAP inference (*maximum a posteriori* inference), in which we are interested in the state of a subset of variables which do not necessarily include all those without evidence.

MPE inference in MLN

- MPE inference in MLN reduces to finding the truth assignment for variables (i.e. nodes) without evidence maximizing the weighted sum of satisfied clauses (i.e. features)
- The problem can be addressed with any weighted satisfiability solver
- MaxWalkSAT has been successfully used for MPE inference in MLN.

Description

- Weighted version of WalkSAT
- Stochastic local search algorithm:
 - 1 Pick an unsatisfied clause at random
 - 2 Flip the truth value of an atom in the clause
- The atom to flip is chosen in one of two possible ways with a certain probability:
 - randomly
 - in order to maximize the weighted sum of the clauses satisfied with the flip
- The stochastic behaviour (hopefully) allows to escape local minima

MaxWalkSAT pseudocode

```
1: procedure MAXWALKSAT(weighted_clauses, max_flips, max_tries, target, p)
2:   vars  $\leftarrow$  variables in weighted_clauses
3:   for i  $\leftarrow$  1 to max_tries do
4:     soln  $\leftarrow$  a random truth assignment to vars
5:     cost  $\leftarrow$  sum of weights of unsatisfied clauses in soln
6:     for j  $\leftarrow$  1 to max_flips do
7:       if cost  $\leq$  target then
8:         return "Success, solution is", soln
9:       end if
10:      c  $\leftarrow$  a randomly chosen unsatisfied clause
11:      if Uniform(0,1) < p then
12:        vf  $\leftarrow$  a randomly chosen variable from c
13:      else
14:        for all variable v in c do
15:          compute DeltaCost(v)
16:        end for
17:        vf  $\leftarrow$  v with lowest DeltaCost(v)
18:      end if
19:      soln  $\leftarrow$  soln with vf flipped
20:      cost  $\leftarrow$  cost + DeltaCost(vf)
21:    end for
22:  end for
23:  return "Failure, best assignment is", best soln found
24: end procedure
```

Ingredients

- *target* is the maximum cost considered acceptable for a solution
- *max_tries* is the number of walk restarts
- *max_flips* is the number of flips in a single walk
- *p* is the probability of flipping a random variable
- $\text{Uniform}(0,1)$ picks a number uniformly at random from $[0,1]$
- $\text{DeltaCost}(v)$ computes the change in cost obtained by flipping variable v in the current solution

Marginal and conditional probabilities

- Another basic inference task is that of computing the marginal probability that a formula holds, possibly given some evidence on the truth value of other formulas
- Exact inference in generic MLN is intractable (as it is for the generic MN obtained by the grounding)
- MCMC sampling techniques have been used as an approximate alternative

Constructing the ground MN

- In order to perform a specific inference task, it is not necessary in general to ground the whole network, as parts of it could have no influence on the computation of the desired probability
- Grounding only the needed part of the network can allow significant savings both in memory and in time to run the inference

Partial grounding: intuition

- A standard inference task is that of computing the probability that F_1 holds given that F_2 does.
- We will focus on the common simple case in which F_1, F_2 are conjunctions of ground literals:
 - 1 All atoms in F_1 are added to the network one after the other
 - 2 If an atom is also in F_2 (has evidence), nothing more is needed for it
 - 3 Otherwise, its Markov blanket is added, and each atom in the blanket is checked in the same way

Partial grounding: pseudocode

```
1: procedure CONSTRUCTNETWORK( $F_1, F_2, L, C$ )  
   inputs:  
    $F_1$  a set of query ground atoms  
    $F_2$  a set of evidence ground atoms  
    $L$  a Markov Logic Network  
    $C$  a set of constants  
   output:  $M$  a ground Markov Network  
   calls:  $MB(q)$  the Markov blanket of  $q$  in  $M_{L,C}$   
2:    $G \leftarrow F_1$   
3:   while  $F_1 \neq \emptyset$  do  
4:     for all  $q \in F_1$  do  
5:       if  $q \notin F_2$  then  
6:          $F_1 \leftarrow F_1 \cup (MB(q) \setminus G)$   
7:          $G \leftarrow G \cup MB(q)$   
8:       end if  
9:        $F_1 \leftarrow F_1 \setminus \{q\}$   
10:    end for  
11:  end while  
12:  return  $M$  the ground MN composed of all nodes in  $G$  and all arcs between  
    them in  $M_{L,C}$ , with features and weights of the corresponding cliques  
13: end procedure
```

Inference

Gibbs sampling

- Inference in the partial ground network is done by Gibbs sampling.
- The basic step consists of sampling a ground atom given its Markov blanket
- The probability of X_I given that its Markov blanket has state $\mathbf{B}_I = \mathbf{b}_I$ is $p(X_I = x_I | \mathbf{B}_I = \mathbf{b}_I) =$

$$\frac{\exp \sum_{f_i \in F_I} w_i f_i(X_I = x_I, \mathbf{B}_I = \mathbf{b}_I)}{\exp \sum_{f_i \in F_I} w_i f_i(X_I = 0, \mathbf{B}_I = \mathbf{b}_I) + \exp \sum_{f_i \in F_I} w_i f_i(X_I = 1, \mathbf{B}_I = \mathbf{b}_I)}$$

where:

- F_I is the set of ground formulas containing X_I
- $f_i(X_I = x_I, \mathbf{B}_I = \mathbf{b}_I)$ is the truth value of the i th formula when $X_I = x_I$ and $\mathbf{B}_I = \mathbf{b}_I$
- The probability of the conjunction of literals is the fraction of samples (at chain convergence) in which all literals are true

Multimodal distributions

- As the distribution is likely to have many modes, multiple independently initialized chains are run
- Efficiency in modeling the multimodal distribution can be obtained starting each chain from a mode reached using MaxWalkSAT

Handling hard constraints

- Hard constraints break the space of possible worlds into separate regions
- This violate the MCMC assumption of reachability
- Very strong constraints create areas of very low probability difficult to traverse
- The problem can be addressed by *slice sampling* MCMC, a technique aimed at sampling from slices of the distribution with a frequency proportional to the probability of the slice

Maximum likelihood parameter estimation

- Parameter estimation amounts at learning weights of formulas
- We can learn weights from training examples as possible worlds.
- Let's consider a single possible world as training example, made of:
 - a set of constants \mathcal{C} defining a specific MN from the MLN
 - a truth value for each ground atom in the resulting MN
- We usually make a *closed world* assumption, where we only specify the true ground atoms, while all others are assumed to be false.
- As all groundings of the same formula will share the same weight, learning can be also done on a single possible world

Maximum likelihood parameter estimation

- Weights of formulas can be learned maximizing the likelihood of the possible world:

$$w^{\max} = \operatorname{argmax}_w p_w(x) = \operatorname{argmax}_w \frac{1}{Z} \exp \left(\sum_{i=1}^F w_i n_i(x) \right)$$

- As usual we will equivalently maximize the log-likelihood:

$$\log(p_w(x)) = \sum_{i=1}^F w_i n_i(x) - \log(Z)$$

Priors

- In order to combat overfitting Gaussian priors can be added to the weights as usual (see CRF)

Maximum likelihood parameter estimation

- The gradient of the log-likelihood wrt weights becomes:

$$\frac{\partial}{\partial w_i} \log p_w(x) = n_i(x) - \sum_{x'} p_w(x') n_i(x')$$

where the sum is over all possible worlds x' , i.e. all possible truth assignments for ground atoms in the MN

- Note that $p_w(x')$ is computed using the current parameter values w
- The i -th component of the gradient is the difference between number of true grounding of the i -th formula, and its expectation according to the current model

Entity resolution

- Determine which observations (e.g. noun phrases in texts) correspond to the same real-world object
- Typically addressed creating feature vectors for pairs of occurrences, and training a classifier to predict whether they match
- The pairwise approach doesn't model information on multiple related objects (e.g. if two bibliographic entries correspond to the same paper, the authors are also the same)
- Some implications hold only with a certain probability (e.g. if two authors in two bibliographic entries are the same, the entries are more likely to refer to the same paper)

MLN for entity resolution

- MLN can be used to address entity resolution tasks by:
 - not assuming that distinct names correspond to distinct objects
 - adding an equality predicate and its axioms: reflexivity, symmetry, transitivity
- Implications related to the equality predicate can be:
 - grounding of a predicate with equal constants have same truth value
 - constants appearing in a ground predicate with equal constants are equal (i.e. the “same paper \rightarrow same author” implication, which holds only probabilistically in general)

MLN for entity resolution

- Weights for different instances of such axioms can be learned from data
- Inference is performed adding evidence on entity properties and relations, and querying for equality atoms
- The network performs *collective* entity resolution, as the most probable resolution for all entities is jointly produced

Entity resolution

Entity resolution in citation databases

- Each citation has: author, title, venue fields.
- Citation to field relations:

```
Author(bib, author)  Title(bib, title)  
Venue(bib, venue)
```

- field content relations:

```
HasWord(author, word)  HasWord(title, word)  
HasWord(venue, word)
```

- equivalence relations:

```
SameAuthor(author1, author2)  
SameTitle(title1, title2)  
SameVenue(venue1, venue2)  
SameBib(bib1, bib2)
```

Entity resolution

Same words imply same entity

- E.g.:

$$\text{Title}(b1, t1) \wedge \text{Title}(b2, t2) \wedge \text{HasWord}(t1, +w) \\ \wedge \text{HasWord}(t2, +w) \Rightarrow \text{SameBib}(b1, b2)$$

- here the '+' operator is a template: a rule is generated for each constant of the appropriate type (i.e. words)
- a separate weight is learned for separate words (e.g. stopwords like articles or prepositions are probably less informative than other words)

transitivity

- E.g.:

$$\text{SameBib}(b1, b2) \wedge \text{SameBib}(b2, b3) \Rightarrow \text{SameBib}(b1, b3)$$

transitivity across entities

- E.g.:

$$\text{Author}(b1, a1) \wedge \text{Author}(b2, a2) \wedge \text{SameBib}(b1, b2) \\ \Rightarrow \text{SameAuthor}(a1, a2)$$
$$\text{Author}(b1, a1) \wedge \text{Author}(b2, a2) \wedge \\ \text{SameAuthor}(a1, a2) \Rightarrow \text{Samebib}(b1, b2)$$

- The second rule is not a valid logic rule, but holds probabilistically (citations with same authors are more likely to be the same)

References

- Domingos, Pedro and Kok, Stanley and Lowd, Daniel and Poon, Hoifung and Richardson, Matthew and Singla, Parag (2007). *Markov Logic*. In L. De Raedt, P. Frasconi, K. Kersting and S. Muggleton (eds.), Probabilistic Inductive Logic Programming. New York: Springer.

Software

- The open source Alchemy system provides an implementation of MLN, with example networks for a number of tasks:
 - <http://alchemy.cs.washington.edu/>